

Language Detection Using NLP & ML

Mini Project –
Advanced Machine Learning (23CSE514)
B.Tech CSE – AI

Introduction

- **Language detection = identifying the language of any text**
- **Used in translation systems, NLP pipelines, chatbots, search engines**
- **This project detects 17 different languages using Machine Learning Model**
- **Uses TF-IDF + Logistic Regression for high accuracy**

Prior Study (Literature Review)

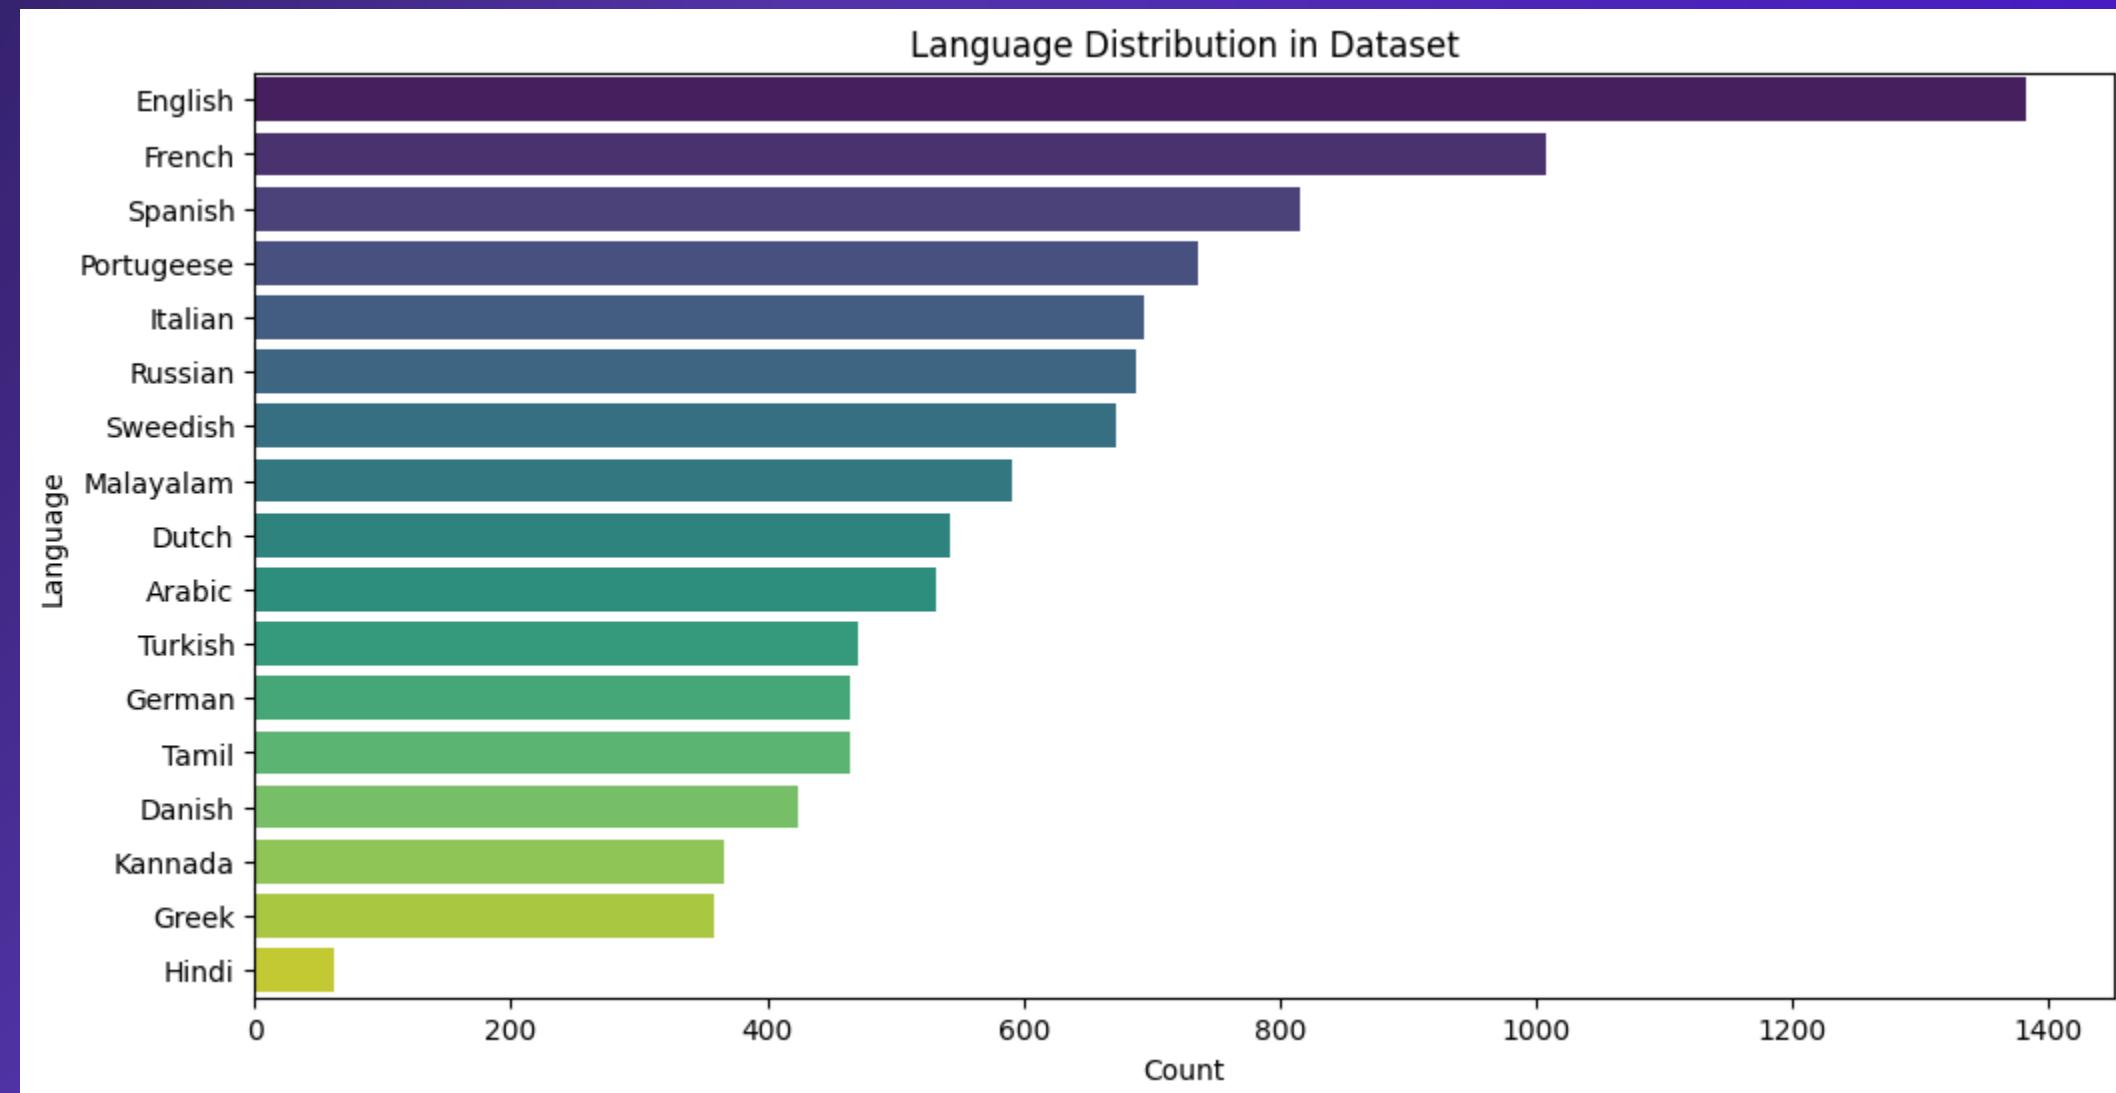
- Traditional rule-based approaches: dictionaries, stopwords (not accurate)
- N-gram models by Cavnar & Trenkle (1994) → high accuracy
- ML models used: Naïve Bayes, Logistic Regression, SVM
- Deep learning exists but heavy for small datasets
- TF-IDF + LR remains most practical and effective for language identification

Problem Statement

- Need automatic language identification from user input text
- Multilingual datasets contain multiple scripts (Latin, Arabic, Devanagari, Dravidian)
- Dataset imbalance leads to bias in predictions
- Goal: Build a robust, accurate, script-independent language detection model

Dataset

- Dataset: Language Detection.csv
- Contains 17 languages
- Around 10,267 samples originally
- Highly imbalanced:
 - English > 1000 samples
 - Hindi ~ 63 samples
- Balancing required to avoid bias



Methodology

- Data Cleaning:
 - Remove nulls
 - Remove duplicates
 - Convert text to lowercase
- Feature extraction: Character-level TF-IDF (2–4 n-grams)
- Balancing: Random Oversampling
- Train-test split: 80% training, 20% testing
- Model: Logistic Regression (multiclass)

How does the model detect a language?

Step 1 → Input comes from user

Example:

“hello”

Step 2 → Convert text into character n-grams

The sentence becomes many pieces like:

For the word "hello":

2-gram (bigram): he, el, ll, lo

3-gram (trigram): hel, ell, llo

4-gram: hell, ello

So, an n-gram helps the model learn the patterns of characters that are unique to each language.

How does the model detect a language?

Step 3 → TF-IDF gives weights to each n-gram

TF-IDF assigns:

High weight → if a character pattern is important/unique

Low weight → if it is common in many languages

This converts text → **numeric vector**.

Step 4 → Logistic Regression uses learned patterns

During training, the model saw thousands of samples per language and learned:

Spanish texts = have patterns like "ξ", "ó", "está", "que"

Hindi = “क”, “न”, “द” etc.

Malayalam = “മത്ത്”, “മെ”, “ഒറ്റ്”

Arabic = “ك”, “م”, “ال”

English = “the”, “ing”, “er”

It learns **character fingerprints**.

How does the model detect a language?

Step 5 → Model predicts the language

Example output:

Spanish: **0.98**

Portuguese: 0.01

Italian: 0.01

So, the predicted language = English.

Code Implementation (Training)

- **TF-IDF vectorizer for character n-grams**
- **RandomOverSampler() for balancing**
- **LogisticRegression(max_iter=2000) for stable training**
- **Save model using joblib**
- **Prediction function defined for inference**

Code Implementation (Training)

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
from imblearn.over_sampling import RandomOverSampler

df = pd.read_csv("/content/Language Detection.csv")

df = df.drop_duplicates(subset=["Text"])
df = df.dropna()

print("Dataset size before balancing:", len(df))

X = df["Text"]
y = df["Language"]

vectorizer = TfidfVectorizer(analyzer="char", ngram_range=(2,4))
X_vec = vectorizer.fit_transform(X)

oversample = RandomOverSampler()
X_bal, y_bal = oversample.fit_resample(X_vec, y)

print("Dataset size after balancing:", len(y_bal))

X_train, X_test, y_train, y_test = train_test_split(
    X_bal, y_bal, test_size=0.2, random_state=42
)
```

```
model = LogisticRegression(max_iter=2000)
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n")
print(classification_report(y_test, y_pred))

def predict_language(text: str) -> str:
    x = vectorizer.transform([text])
    return model.predict(x)[0]

def interactive_console():
    print("\n== Interactive Language Detection (type 'q' to quit) ==")
    while True:
        user_text = input("\nEnter text: ")
        if user_text.strip().lower() in ["q", "quit", "exit"]:
            print("Exiting...")
            break
        if not user_text.strip():
            print("Please enter some text.")
            continue

        lang = predict_language(user_text)
        print("Predicted Language:", lang)

try:
    import ipywidgets as widgets
```

```
try:
    import ipywidgets as widgets
    from IPython.display import display, Markdown

    text_box = widgets.Textarea(
        value="Type some text here...", 
        placeholder="Enter text to detect language",
        description="Input:",
        layout=widgets.Layout(width="100%", height="120px")
    )

    button = widgets.Button(
        description="Detect Language",
        button_style="", # 'success', 'info', 'warning', 'danger' or ''
        tooltip="Click to detect language"
    )

    output = widgets.Output()

    def on_button_click(b):
        with output:
            output.clear_output()
            user_text = text_box.value.strip()
            if not user_text:
                print("Please enter some text.")
                return
            lang = predict_language(user_text)
            display(Markdown(f"**Predicted Language:** {lang}"))

    button.on_click(on_button_click)

    display(Markdown("## Language Detection"))
    display(text_box, button, output)

except ImportError:
    print("ipywidgets not available - only console mode can be used.")
```

Model Evaluation

```
Dataset size before balancing: 10267  
Dataset size after balancing: 23494  
Accuracy: 0.9970206426899341
```

Classification Report:

	precision	recall	f1-score	support
Arabic	1.00	1.00	1.00	248
Danish	1.00	1.00	1.00	264
Dutch	1.00	0.99	1.00	297
English	0.99	1.00	0.99	250
French	1.00	1.00	1.00	308
German	0.99	1.00	0.99	291
Greek	1.00	1.00	1.00	266
Hindi	1.00	1.00	1.00	278
Italian	0.99	0.99	0.99	279
Kannada	1.00	1.00	1.00	267
Malayalam	1.00	1.00	1.00	271
Portugeese	0.99	0.99	0.99	291
Russian	1.00	1.00	1.00	288
Spanish	0.99	0.99	0.99	282
Sweedish	1.00	1.00	1.00	281
Tamil	1.00	1.00	1.00	259
Turkish	1.00	1.00	1.00	279
accuracy			1.00	4699
macro avg	1.00	1.00	1.00	4699
weighted avg	1.00	1.00	1.00	4699

Code Implementation: Output

Language Detection

Input: छपाई और अक्षर योजन उद्योग का एक साधारण डमी पाठ है.

Detect Language

Predicted Language: Hindi

Language Detection

Input: Lorem Ipsum es un tipo de texto de marcador de posición que se usa comúnmente en las industrias del

Detect Language

Predicted Language: Spanish

Language Detection

Input: Lorem Ipsum, bir sayfadaki boşluğu doldurmak ve nihai içeriğin nasıl görüneceğine dair bir izlenim

Detect Language

Predicted Language: Turkish

Code Implementation: Output

Language Detection

Input: പ്രസിലീകരണ, ഗ്രാഫിക് ഡിസൈൻ മേഖലകളിൽ ഉപയോഗിക്കുന്ന ഒരു പദ്ധതി

Detect Language

Predicted Language: Malayalam

Language Detection

Input: ലോറമ് ഇപ്സെൻസ് എന്നുള്ള ഒരു റീഥിയ ഫ്ലോറൽ പ്രിൻടിംഗ്

Detect Language

Predicted Language: Kannada

Code Implementation (Flask App)

- Load saved `lang_model.pkl` and `vectorizer.pkl`
- Flask routes:
 - `/` → detect language
 - `/performance` → show model accuracy & F1-scores
- UI: text input box → predicted language
- Confidence score displayed using `predict_proba()`
- Modern UI with HTML/CSS

Code Implementation (Flask App)

The image shows a screenshot of a Flask web application interface. The main title is "Detect Language from Text" with a subtitle "Uses TF-IDF + Logistic Regression trained on 17 languages". Below this is a text input field labeled "Enter text" containing the Hindi sentence "छपाई और अक्षर योजन उद्योग का एक साधारण डमी पाठ है". A large blue button labeled "DETECT LANGUAGE" is centered below the text input. The result section displays "Predicted Language: Hindi" with a "High confidence" badge. A green progress bar indicates "Model confidence" at 36.32%. A table titled "Top predicted languages with probabilities" lists the following data:

Language	Probability
Hindi	36.32%
Arabic	5.73%
Russian	5.4%
Turkish	5.16%
Greek	4.77%

Code Implementation (Flask App)

Detect Language from Text

Uses TF-IDF + Logistic Regression trained on 17 languages

Enter text

hello, testing the model

DETECT LANGUAGE

Predicted Language: English High confidence

Model confidence 75.07%

Top predicted languages with probabilities

Language	Probability
English	75.07%
Danish	3.29%
Spanish	2.7%
Italian	2.59%
Dutch	2.17%

Detect Language from Text

Uses TF-IDF + Logistic Regression trained on 17 languages

Enter text

ஹோஹம் இப்சம் என்பது ஒரு பக்கத்தில் உள்ள இடத்தை நிரப்பவும்

DETECT LANGUAGE

Predicted Language: Tamil High confidence

Model confidence 97.62%

Top predicted languages with probabilities

Language	Probability
Tamil	97.62%
Arabic	0.2%
Russian	0.2%
Turkish	0.19%
Greek	0.18%

Code Implementation (Flask App)

Model Performance				
Evaluation on a held-out test split of the Language Detection dataset.				
Accuracy 0.999	Macro F1-score 1.0	Weighted F1-score 1.0	Macro Precision / Recall 1.0 / 1.0	
Per-class metrics				
Language	Precision	Recall	F1-score	Support
Arabic	1.0	1.0	1.0	106
Danish	1.0	1.0	1.0	85
Dutch	1.0	1.0	1.0	108
English	1.0	1.0	1.0	277
French	1.0	1.0	1.0	201
German	1.0	1.0	1.0	93
Greek	1.0	1.0	1.0	72
Hindi	1.0	1.0	1.0	12
Italian	0.99	1.0	0.99	139
Kannada	1.0	1.0	1.0	73
Malayalam	1.0	1.0	1.0	118
Portugeese	1.0	0.99	1.0	147
Russian	1.0	1.0	1.0	138
Spanish	1.0	0.99	1.0	163
Sweedish	1.0	1.0	1.0	135
Tamil	1.0	1.0	1.0	93
Turkish	1.0	1.0	1.0	94

Results

- **Accuracy: 99.9%**
- **Macro F1-score: 1.0**
- **Weighted F1-score: 1.0**
- **Perfect per-language results**
- **Balanced dataset fixed model bias**
- **Flask app predicts accurately in real-time**

Summary

- Complete NLP + ML pipeline implemented
- Very high performance achieved
- Character-level n-grams ideal for language detection
- Balanced dataset ensured fairness
- Deployed using Flask for real-world usage

References

- **Cavnar & Trenkle (1994)** – N-gram text categorization
- **Manning et al.** – *Introduction to Information Retrieval*
- **scikit-learn documentation**
- **imbalanced-learn documentation**
- **Flask documentation**

Thank You